

Homework 2: Dispatcher/Worker Model Solution Description Student IDs: 207477753, 208076919

1. Introduction This project implements a multi-threaded Dispatcher/Worker system in C for Linux. It processes a command file where specific "dispatcher" commands are executed serially, while "worker" commands are offloaded to a thread pool via a shared job queue. The solution prioritizes data integrity and concurrency control using POSIX threads (pthreads) and synchronization primitives.

2. High-Level Design The system comprises two main components:

- **Dispatcher (Main Thread):** Parses the input command file, initializes the system, and manages the lifecycle of the worker threads. It handles dispatcher_msleep and dispatcher_wait commands directly.
- **Worker Pool:** A fixed number of threads (created via pthread_create) that consume jobs from a shared queue. Each worker processes complex command lines (containing msleep, increment, decrement, repeat) independently.

3. Data Structures

- **Job Queue (job_queue_t):** A singly linked list is used to manage pending jobs. Each node (job_t) contains the command string and the timestamp of when it was read (read_time_ms), allowing for accurate turnaround time calculations.
- **Counter Files:** Integers are stored in persistent text files (countXX.txt).
- **Logging:** Trace data is written to dispatcher.txt and threadXX.txt to track execution flow.

4. Synchronization & Concurrency To ensure thread safety and maximize parallelism, the following mechanisms are used:

- **Queue Protection:** A global mutex (queue_mutex) protects all operations on the work_queue (enqueue/dequeue).
+1
- **Condition Variables:**
 - queue_not_empty: Signals sleeping workers when a new job is available, preventing busy waiting.
+2
 - all_jobs_finished: Signals the dispatcher when the queue is empty and all active workers are idle, enabling the dispatcher_wait command.
+1
- **Fine-Grained File Locking:** An array of mutexes (file_mutexes[MAX_COUNTERS]) is used to protect counter files. This allows multiple workers to update *different* counters

simultaneously without blocking each other, while ensuring atomic updates for the *same* counter.

+1

- **Statistics Protection:** A stats_mutex ensures that updates to global statistics (sum, min, max turnaround times) are atomic.

+1

5. Implementation Details

- **Command Parsing:**
 - The dispatcher uses strtok to parse the main command file.
 - Worker threads use the reentrant function strtok_r to parse job lines safely without interfering with other threads' string processing.
- **File Operations:** Counters are updated using a read-modify-write cycle: fopen -> fscanf -> rewind -> fprintf -> ftruncate -> fclose. This ensures the file content remains clean.
+1
- **Memory Management:** The system dynamically allocates memory for job nodes and the worker thread pool. A comprehensive cleanup routine runs at the end of the dispatcher function to join all threads, free allocated arrays, and destroy all mutexes and condition variables